

```
In [1]: #BERT code is from https://github.com/google-research/bert/blob/master/predicting\_movie\_reviews\_with\_bert\_on\_tf\_hub.ipynb
```

```
In [ ]: !pip install bert-tensorflow
```

```
In [2]: import pandas as pd
from datetime import datetime
import tensorflow as tf
import tensorflow_hub as hub
from sklearn.model_selection import train_test_split

import bert
from bert import run_classifier
from bert import optimization
from bert import tokenization
```

WARNING:tensorflow:From C:\Users\alex\Anaconda3\envs\CSML1010_OnlineSession2\lib\site-packages\bert\optimization.py:87: The name tf.train.Optimizer is deprecated. Please use tf.compat.v1.train.Optimizer instead.


```
In [3]: df_tweets_cleaned = pd.read_csv('../data/Tweets_cleaned.csv')
```

```
In [4]: df_tweets_cleaned
```

```
Out[4]:
```

	tweet_id	airline_sentiment	airline_sentiment_confidence	negativereason	negativereason_confidence	airline
0	570306133677760513	neutral	1.0000	NaN	NaN	Virgin America
1	570301130888122368	positive	0.3486	NaN	0.0000	Virgin America
2	570301083672813571	neutral	0.6837	NaN	NaN	Virgin America
3	570301031407624196	negative	1.0000	Bad Flight	0.7033	Virgin America
4	570300817074462722	negative	1.0000	Can't Tell	1.0000	Virgin America
...
14635	569587686496825344	positive	0.3487	NaN	0.0000	American Airlines
14636	569587371693355008	negative	1.0000	Customer Service Issue	1.0000	American Airlines
14637	569587242672398336	neutral	1.0000	NaN	NaN	American Airlines
14638	569587188687634433	negative	1.0000	Customer Service Issue	0.6659	American Airlines
14639	569587140490866689	neutral	0.6771	NaN	0.0000	American Airlines

14640 rows × 7 columns



```
In [5]: #df_tweets_bert = df_tweets_cleaned[['airline_sentiment', 'text_cleaned']]
df_tweets_bert = df_tweets_cleaned[['airline_sentiment', 'text_list_no_stop_words']]
```

```
In [6]: df_tweets_bert
```

```
Out[6]:
```

	airline_sentiment	text_list_no_stop_words
0	neutral	said
1	positive	plus added commercials experience tacky
2	neutral	today mean need trip
3	negative	aggressive blast obnoxious entertainment guest...
4	negative	big bad thing
...
14635	positive	thank got different flight chicago
14636	negative	leaving minutes late flight warnings communica...
14637	neutral	bring american airlines
14638	negative	money change flight answer phones suggestions ...
14639	neutral	people need know seats flight standby people f...

14640 rows × 2 columns

```
In [7]: df_tweets_bert['binary_response_variable'] = False

df_tweets_bert.loc[df_tweets_cleaned.airline_sentiment == 'neutral', 'binary_response_variable'] = False
df_tweets_bert.loc[df_tweets_cleaned.airline_sentiment == 'positive', 'binary_response_variable'] = False
df_tweets_bert.loc[df_tweets_cleaned.airline_sentiment == 'negative', 'binary_response_variable'] = True
```

C:\Users\alex\Anaconda3\envs\CSML1010_OnlineSession2\lib\site-packages\ipykernel_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

"""Entry point for launching an IPython kernel.

C:\Users\alex\Anaconda3\envs\CSML1010_OnlineSession2\lib\site-packages\pandas\core\indexing.py:966: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

self.obj[item] = s

```
In [8]: df_tweets_bert.binary_response_variable = df_tweets_bert.binary_response_variable.astype(int)
```

C:\Users\alex\Anaconda3\envs\CSML1010_OnlineSession2\lib\site-packages\pandas\core\generic.py:5303: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

self[name] = value

```
In [9]: df_tweets_bert = df_tweets_bert[['text_list_no_stop_words', 'binary_response_variable']]
df_tweets_bert
```

Out[9]:

	text_list_no_stop_words	binary_response_variable
0	said	0
1	plus added commercials experience tacky	0
2	today mean need trip	0
3	aggressive blast obnoxious entertainment guest...	1
4	big bad thing	1
...
14635	thank got different flight chicago	0
14636	leaving minutes late flight warnings communica...	1
14637	bring american airlines	0
14638	money change flight answer phones suggestions ...	1
14639	people need know seats flight standby people f...	0

14640 rows × 2 columns

```
In [10]: X = df_tweets_bert.loc[:, df_tweets_bert.columns != 'binary_response_variable']
Y = df_tweets_bert.loc[:, df_tweets_bert.columns == 'binary_response_variable']

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.30, random_state=1, stratify=Y)
```

```
In [11]: X_train
```

Out[11]:

	text_list_no_stop_words
4771	arrangements reimburse rental
12454	apologizing rude sales reps failure offer trit...
3840	year old flying tokyo vacation bad knees happens
14176	baggage lost flight cancelled flightled accomm...
1219	passengers rerouted match intended arrival tim...
...	...
7233	yeah aware original message
3855	free upgrade problems reservation mkwlkr
6339	hey guys honest customers unlike
11606	jfk baggage office open help book cancelled fl...
3899	pri boarding active military uniform travel un...

10248 rows × 1 columns

```
In [12]: XY_train = pd.concat([X_train, Y_train], axis=1)
XY_test = pd.concat([X_test, Y_test], axis=1)
```

```
In [13]: XY_train = XY_train.dropna()
XY_test = XY_test.dropna()
```

```
In [14]: print(XY_train.isnull().values.any())
print(XY_test.isnull().values.any())
```

False
False

```
In [15]: DATA_COLUMN = 'text_list_no_stop_words'
        LABEL_COLUMN = 'binary_response_variable'

        # label_list is the list of labels. In this situation they are 0, 1
        label_list = [0, 1]
```

```
In [16]: XY_train
```

```
Out[16]:
```

	text_list_no_stop_words	binary_response_variable
4771	arrangements reimburse rental	0
12454	apologizing rude sales reps failure offer trit...	1
3840	year old flying tokyo vacation bad knees happens	1
14176	baggage lost flight cancelled flightled accomm...	1
1219	passengers rerouted match intended arrival tim...	1
...
7233	yeah aware original message	1
3855	free upgrade problems reservation mkwlkr	1
6339	hey guys honest customers unlike	0
11606	jfk baggage office open help book cancelled fl...	1
3899	pri boarding active military uniform travel un...	0

10190 rows × 2 columns

Data Preprocessing

```
In [17]: # Use the InputExample class from BERT's run_classifier code to create examples from the data
        train_InputExamples = XY_train.apply(lambda x: bert.run_classifier.InputExample(guid=None, # Globally
        unique ID for bookkeeping, unused in this example
        text_a = x[DATA_COLUMN],
        text_b = None,
        label = x[LABEL_COLUMN]), axis = 1)

        test_InputExamples = XY_test.apply(lambda x: bert.run_classifier.InputExample(guid=None,
        text_a = x[DATA_COLUMN],
        text_b = None,
        label = x[LABEL_COLUMN]), axis = 1)
```

```
In [18]: train_InputExamples
```

```
Out[18]: 4771    <bert.run_classifier.InputExample object at 0x...
        12454    <bert.run_classifier.InputExample object at 0x...
        3840    <bert.run_classifier.InputExample object at 0x...
        14176    <bert.run_classifier.InputExample object at 0x...
        1219    <bert.run_classifier.InputExample object at 0x...
        ...
        7233    <bert.run_classifier.InputExample object at 0x...
        3855    <bert.run_classifier.InputExample object at 0x...
        6339    <bert.run_classifier.InputExample object at 0x...
        11606    <bert.run_classifier.InputExample object at 0x...
        3899    <bert.run_classifier.InputExample object at 0x...
        Length: 10190, dtype: object
```

```
In [19]: # This is a path to an uncased (all lowercase) version of BERT
BERT_MODEL_HUB = "https://tfhub.dev/google/bert_uncased_L-12_H-768_A-12/1"

def create_tokenizer_from_hub_module():
    """Get the vocab file and casing info from the Hub module."""
    with tf.Graph().as_default():
        bert_module = hub.Module(BERT_MODEL_HUB)
        tokenization_info = bert_module(signature="tokenization_info", as_dict=True)
        with tf.Session() as sess:
            vocab_file, do_lower_case = sess.run([tokenization_info["vocab_file"],
                                                  tokenization_info["do_lower_case"]])

    return bert.tokenization.FullTokenizer(
        vocab_file=vocab_file, do_lower_case=do_lower_case)

tokenizer = create_tokenizer_from_hub_module()
```

INFO:tensorflow:Saver not created because there are no variables in the graph to restore

INFO:tensorflow:Saver not created because there are no variables in the graph to restore

WARNING:tensorflow:From C:\Users\alex\Anaconda3\envs\CSML1010_OnlineSession2\lib\site-packages\bert\tokenization.py:125: The name tf.gfile.GFile is deprecated. Please use tf.io.gfile.GFile instead.

WARNING:tensorflow:From C:\Users\alex\Anaconda3\envs\CSML1010_OnlineSession2\lib\site-packages\bert\tokenization.py:125: The name tf.gfile.GFile is deprecated. Please use tf.io.gfile.GFile instead.

```
In [20]: # We'll set sequences to be at most 128 tokens Long.  
MAX_SEQ_LENGTH = 128  
# Convert our train and test features to InputFeatures that BERT understands.  
train_features = bert.run_classifier.convert_examples_to_features(train_InputExamples, label_list, MAX_SEQ_LENGTH, tokenizer)  
test_features = bert.run_classifier.convert_examples_to_features(test_InputExamples, label_list, MAX_SEQ_LENGTH, tokenizer)
```

[illegible]

[illegible]

[illegible]

[illegible]


```
In [21]: train_InputExamples
```

```
Out[21]: 4771      <bert.run_classifier.InputExample object at 0x...  
12454      <bert.run_classifier.InputExample object at 0x...  
3840      <bert.run_classifier.InputExample object at 0x...  
14176      <bert.run_classifier.InputExample object at 0x...  
1219      <bert.run_classifier.InputExample object at 0x...  
          ...  
7233      <bert.run_classifier.InputExample object at 0x...  
3855      <bert.run_classifier.InputExample object at 0x...  
6339      <bert.run_classifier.InputExample object at 0x...  
11606      <bert.run_classifier.InputExample object at 0x...  
3899      <bert.run_classifier.InputExample object at 0x...  
Length: 10190, dtype: object
```

Model

```

In [22]: def create_model(is_predicting, input_ids, input_mask, segment_ids, labels,
                        num_labels):
    """Creates a classification model."""

    bert_module = hub.Module(
        BERT_MODEL_HUB,
        trainable=True)
    bert_inputs = dict(
        input_ids=input_ids,
        input_mask=input_mask,
        segment_ids=segment_ids)
    bert_outputs = bert_module(
        inputs=bert_inputs,
        signature="tokens",
        as_dict=True)

    # Use "pooled_output" for classification tasks on an entire sentence.
    # Use "sequence_outputs" for token-level output.
    output_layer = bert_outputs["pooled_output"]

    hidden_size = output_layer.shape[-1].value

    # Create our own layer to tune for politeness data.
    output_weights = tf.get_variable(
        "output_weights", [num_labels, hidden_size],
        initializer=tf.truncated_normal_initializer(stddev=0.02))

    output_bias = tf.get_variable(
        "output_bias", [num_labels], initializer=tf.zeros_initializer())

    with tf.variable_scope("loss"):

        # Dropout helps prevent overfitting
        output_layer = tf.nn.dropout(output_layer, keep_prob=0.9)

        logits = tf.matmul(output_layer, output_weights, transpose_b=True)
        logits = tf.nn.bias_add(logits, output_bias)
        log_probs = tf.nn.log_softmax(logits, axis=-1)

        # Convert labels into one-hot encoding
        one_hot_labels = tf.one_hot(labels, depth=num_labels, dtype=tf.float32)

        predicted_labels = tf.squeeze(tf.argmax(log_probs, axis=-1, output_type=tf.int32))
        # If we're predicting, we want predicted labels and the probabilities.
        if is_predicting:
            return (predicted_labels, log_probs)

        # If we're train/eval, compute loss between predicted and actual label
        per_example_loss = -tf.reduce_sum(one_hot_labels * log_probs, axis=-1)
        loss = tf.reduce_mean(per_example_loss)
        return (loss, predicted_labels, log_probs)

```

```

In [23]: # model_fn_builder actually creates our model function
# using the passed parameters for num_labels, learning_rate, etc.
def model_fn_builder(num_labels, learning_rate, num_train_steps,
                    num_warmup_steps):
    """Returns `model_fn` closure for TPUEstimator."""
    def model_fn(features, labels, mode, params): # pylint: disable=unused-argument
        """The `model_fn` for TPUEstimator."""

        input_ids = features["input_ids"]
        input_mask = features["input_mask"]
        segment_ids = features["segment_ids"]
        label_ids = features["label_ids"]

        is_predicting = (mode == tf.estimator.ModeKeys.PREDICT)

        # TRAIN and EVAL
        if not is_predicting:

            (loss, predicted_labels, log_probs) = create_model(
                is_predicting, input_ids, input_mask, segment_ids, label_ids, num_labels)

            train_op = bert.optimization.create_optimizer(
                loss, learning_rate, num_train_steps, num_warmup_steps, use_tpu=False)

            # Calculate evaluation metrics.
            def metric_fn(label_ids, predicted_labels):
                accuracy = tf.metrics.accuracy(label_ids, predicted_labels)
                f1_score = tf.contrib.metrics.f1_score(
                    label_ids,
                    predicted_labels)
                f2_score = tf.contrib.metrics.f2_score(
                    label_ids,
                    predicted_labels)
            #
            #
            #
            auc = tf.metrics.auc(
                label_ids,
                predicted_labels)
            recall = tf.metrics.recall(
                label_ids,
                predicted_labels)
            precision = tf.metrics.precision(
                label_ids,
                predicted_labels)
            true_pos = tf.metrics.true_positives(
                label_ids,
                predicted_labels)
            true_neg = tf.metrics.true_negatives(
                label_ids,
                predicted_labels)
            false_pos = tf.metrics.false_positives(
                label_ids,
                predicted_labels)
            false_neg = tf.metrics.false_negatives(
                label_ids,
                predicted_labels)
            return {
                "eval_accuracy": accuracy,
                "f1_score": f1_score,
                "f2_score": f2_score,
                "auc": auc,
                "precision": precision,
                "recall": recall,
                "true_positives": true_pos,
                "true_negatives": true_neg,
                "false_positives": false_pos,
                "false_negatives": false_neg
            }
            #

        eval_metrics = metric_fn(label_ids, predicted_labels)

        if mode == tf.estimator.ModeKeys.TRAIN:

```



```

        return tf.estimator.EstimatorSpec(mode=mode,
            loss=loss,
            train_op=train_op)
    else:
        return tf.estimator.EstimatorSpec(mode=mode,
            loss=loss,
            eval_metric_ops=eval_metrics)
    else:
        (predicted_labels, log_probs) = create_model(
            is_predicting, input_ids, input_mask, segment_ids, label_ids, num_labels)

        predictions = {
            'probabilities': log_probs,
            'labels': predicted_labels
        }
        return tf.estimator.EstimatorSpec(mode, predictions=predictions)

# Return the actual model function in the closure
    return model_fn

```

```

In [24]: BATCH_SIZE = 1
        LEARNING_RATE = 2e-5
        NUM_TRAIN_EPOCHS = 3.0

        WARMUP_PROPORTION = 0.1

        SAVE_CHECKPOINTS_STEPS = 500
        SAVE_SUMMARY_STEPS = 100

```

```

In [25]: # Compute # train and warmup steps from batch size
        num_train_steps = int(len(train_features) / BATCH_SIZE * NUM_TRAIN_EPOCHS)
        num_warmup_steps = int(num_train_steps * WARMUP_PROPORTION)

```

```

In [26]: OUTPUT_DIR = 'bert'

# Specify output directory and number of checkpoint steps to save
    run_config = tf.estimator.RunConfig(
        model_dir=OUTPUT_DIR,
        save_summary_steps=SAVE_SUMMARY_STEPS,
        save_checkpoints_steps=SAVE_CHECKPOINTS_STEPS)

```

```
In [27]: model_fn = model_fn_builder(
    num_labels=len(label_list),
    learning_rate=LEARNING_RATE,
    num_train_steps=num_train_steps,
    num_warmup_steps=num_warmup_steps)

estimator = tf.estimator.Estimator(
    model_fn=model_fn,
    config=run_config,
    params={"batch_size": BATCH_SIZE})

INFO:tensorflow:Using config: {'_model_dir': 'bert', '_tf_random_seed': None, '_save_summary_steps':
100, '_save_checkpoints_steps': 500, '_save_checkpoints_secs': None, '_session_config': allow_soft_pl
acement: true
graph_options {
  rewrite_options {
    meta_optimizer_iterations: ONE
  }
}
, '_keep_checkpoint_max': 5, '_keep_checkpoint_every_n_hours': 10000, '_log_step_count_steps': 100,
'_train_distribute': None, '_device_fn': None, '_protocol': None, '_eval_distribute': None, '_experim
ental_distribute': None, '_experimental_max_worker_delay_secs': None, '_session_creation_timeout_sec
s': 7200, '_service': None, '_cluster_spec': <tensorflow.python.training.server_lib.ClusterSpec objec
t at 0x00000272098F33C8>, '_task_type': 'worker', '_task_id': 0, '_global_id_in_cluster': 0, '_maste
r': '', '_evaluation_master': '', '_is_chief': True, '_num_ps_replicas': 0, '_num_worker_replicas':
1}

INFO:tensorflow:Using config: {'_model_dir': 'bert', '_tf_random_seed': None, '_save_summary_steps':
100, '_save_checkpoints_steps': 500, '_save_checkpoints_secs': None, '_session_config': allow_soft_pl
acement: true
graph_options {
  rewrite_options {
    meta_optimizer_iterations: ONE
  }
}
, '_keep_checkpoint_max': 5, '_keep_checkpoint_every_n_hours': 10000, '_log_step_count_steps': 100,
'_train_distribute': None, '_device_fn': None, '_protocol': None, '_eval_distribute': None, '_experim
ental_distribute': None, '_experimental_max_worker_delay_secs': None, '_session_creation_timeout_sec
s': 7200, '_service': None, '_cluster_spec': <tensorflow.python.training.server_lib.ClusterSpec objec
t at 0x00000272098F33C8>, '_task_type': 'worker', '_task_id': 0, '_global_id_in_cluster': 0, '_maste
r': '', '_evaluation_master': '', '_is_chief': True, '_num_ps_replicas': 0, '_num_worker_replicas':
1}
```

```
In [28]: # Create an input function for training. drop_remainder = True for using TPUs.
train_input_fn = bert.run_classifier.input_fn_builder(
    features=train_features,
    seq_length=MAX_SEQ_LENGTH,
    is_training=True,
    drop_remainder=False)
```

```
In [29]: print(f'Beginning Training!')
current_time = datetime.now()
estimator.train(input_fn=train_input_fn, max_steps=num_train_steps)
print("Training took time ", datetime.now() - current_time)
```

```
Beginning Training!
INFO:tensorflow:Skipping training since max_steps has already saved.

INFO:tensorflow:Skipping training since max_steps has already saved.

Training took time  0:00:00.048015
```

```
In [30]: test_input_fn = run_classifier.input_fn_builder(
    features=test_features,
    seq_length=MAX_SEQ_LENGTH,
    is_training=False,
    drop_remainder=False)
```

```
In [43]: #print(f'Beginning predictions!')
#current_time = datetime.now()
estimator.evaluate(input_fn=test_input_fn, steps=None)
#print("Training predictions took this amount of time ", datetime.now() - current_time)
```

INFO:tensorflow:Calling model_fn.

INFO:tensorflow:Calling model_fn.

INFO:tensorflow:Saver not created because there are no variables in the graph to restore

INFO:tensorflow:Saver not created because there are no variables in the graph to restore

C:\Users\alex\Anaconda3\envs\CSML1010_OnlineSession2\lib\site-packages\tensorflow_core\python\framework\indexed_slices.py:424: UserWarning: Converting sparse IndexedSlices to a dense Tensor of unknown shape. This may consume a large amount of memory.

"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

INFO:tensorflow:Done calling model_fn.

INFO:tensorflow:Done calling model_fn.

INFO:tensorflow:Starting evaluation at 2020-05-22T09:02:40Z

INFO:tensorflow:Starting evaluation at 2020-05-22T09:02:40Z

INFO:tensorflow:Graph was finalized.

INFO:tensorflow:Graph was finalized.

INFO:tensorflow:Restoring parameters from bert\model.ckpt-30570

INFO:tensorflow:Restoring parameters from bert\model.ckpt-30570

INFO:tensorflow:Running local_init_op.

INFO:tensorflow:Running local_init_op.

INFO:tensorflow:Done running local_init_op.

INFO:tensorflow:Done running local_init_op.

INFO:tensorflow:Finished evaluation at 2020-05-22-09:03:53

INFO:tensorflow:Finished evaluation at 2020-05-22-09:03:53

INFO:tensorflow:Saving dict for global step 30570: auc = 0.81412077, eval_accuracy = 0.83272314, f1_score = 0.86944085, false_negatives = 311.0, false_positives = 420.0, global_step = 30570, loss = 1.1803308, precision = 0.8528381, recall = 0.8867031, true_negatives = 1205.0, true_positives = 2434.0

INFO:tensorflow:Saving dict for global step 30570: auc = 0.81412077, eval_accuracy = 0.83272314, f1_score = 0.86944085, false_negatives = 311.0, false_positives = 420.0, global_step = 30570, loss = 1.1803308, precision = 0.8528381, recall = 0.8867031, true_negatives = 1205.0, true_positives = 2434.0

INFO:tensorflow:Saving 'checkpoint_path' summary for global step 30570: bert\model.ckpt-30570

INFO:tensorflow:Saving 'checkpoint_path' summary for global step 30570: bert\model.ckpt-30570

```
Out[43]: {'auc': 0.81412077,
'eval_accuracy': 0.83272314,
'f1_score': 0.86944085,
'false_negatives': 311.0,
'false_positives': 420.0,
'loss': 1.1803308,
'precision': 0.8528381,
'recall': 0.8867031,
'true_negatives': 1205.0,
'true_positives': 2434.0,
'global_step': 30570}
```

```
In [37]: def getPrediction(in_sentences):
        labels = [0, 1]
        input_examples = [run_classifier.InputExample(guid="", text_a = x, text_b = None, label = 0) for x, i
        n in zip(in_sentences, range(len(in_sentences)))] # here, "" is just a dummy label
        input_features = run_classifier.convert_examples_to_features(input_examples, label_list, MAX_SEQ_LENGTH, tokenizer)
        predict_input_fn = run_classifier.input_fn_builder(features=input_features, seq_length=MAX_SEQ_LENGTH, is_training=False, drop_remainder=False)
        predictions = estimator.predict(predict_input_fn)
        return [(sentence, prediction['probabilities'], labels[prediction['labels']]) for sentence, prediction in zip(in_sentences, predictions)]
```

```
In [41]: #pred_sentences = [  
#   "My flight to New York was delayedn AGAIN"#,  
#   #"So excited to be on the flight to San Francisco!"  
#]  
  
#predictions = getPrediction(pred_sentences)  
#predictions
```

[illegible]

```

-----
IndexError                                Traceback (most recent call last)
<ipython-input-41-2848935e4500> in <module>
----> 1 predictions = getPrediction(pred_sentences)
      2 predictions

<ipython-input-37-0fd105d73268> in getPrediction(in_sentences)
      5 predict_input_fn = run_classifier.input_fn_builder(features=input_features, seq_length=MAX_
SEQ_LENGTH, is_training=False, drop_remainder=False)
      6 predictions = estimator.predict(predict_input_fn)
----> 7 return [(sentence, prediction['probabilities'], labels[prediction['labels']]) for sentence,
prediction in zip(in_sentences, predictions)]

<ipython-input-37-0fd105d73268> in <listcomp>(.0)
      5 predict_input_fn = run_classifier.input_fn_builder(features=input_features, seq_length=MAX_
SEQ_LENGTH, is_training=False, drop_remainder=False)
      6 predictions = estimator.predict(predict_input_fn)
----> 7 return [(sentence, prediction['probabilities'], labels[prediction['labels']]) for sentence,
prediction in zip(in_sentences, predictions)]

~\Anaconda3\envs\CSML1010_OnlineSession2\lib\site-packages\tensorflow_estimator\python\estimator\esti
mator.py in predict(self, input_fn, predict_keys, hooks, checkpoint_path, yield_single_examples)
    645         yield pred
    646     else:
--> 647         for i in range(self._extract_batch_length(preds_evaluated)):
    648             yield {
    649                 key: value[i]

~\Anaconda3\envs\CSML1010_OnlineSession2\lib\site-packages\tensorflow_estimator\python\estimator\esti
mator.py in _extract_batch_length(self, preds_evaluated)
   1030     for key, value in six.iteritems(preds_evaluated):
   1031         batch_length = batch_length or value.shape[0]
-> 1032     if value.shape[0] != batch_length:
   1033         raise ValueError('Batch length of predictions should be same. %s has '
   1034                             'different batch length than others.' % key)

IndexError: tuple index out of range

```

References

Code for this project is either directly from (with some modification), or inspired by, but not limited to the following sources:

- Predicting Movie Review Sentiment with BERT on TF Hub: https://github.com/google-research/bert/blob/master/predicting_movie_reviews_with_bert_on_tf_hub.ipynb (https://github.com/google-research/bert/blob/master/predicting_movie_reviews_with_bert_on_tf_hub.ipynb)
- Introducing BERT with Tensorflow: <https://www.kaggle.com/sergeykalutsky/introducing-bert-with-tensorflow> (<https://www.kaggle.com/sergeykalutsky/introducing-bert-with-tensorflow>)

In []: